

Gen AI & Classic ML

when to use what

July 2024

Felix Wick

Classic ML vs Deep Learning

text generation

```
done = False
total_reward = 0
while not done:
    state = torch.tensor(state, dtype=torch.float32)
    logits = policy_network(state)
    action = torch.argmax(logits).item()
    next_state, reward, done, _ = env.step(action)
    total_reward += reward
    state = next_state
    total_rewards.append(total_reward)

# Print average reward
print("Average Reward:", np.mean(total_rewards))
```

This code uses PyTorch to implement the policy gradient method (REINFORCE algorithm) to solve the CartPole problem. The policy network is defined as a simple feedforward neural network, and the training loop updates the policy network parameters to maximize the expected reward. Finally, it evaluates the learned policy by running it for 100 episodes and prints the average reward. Adjust hyperparameters and network architecture as needed for better performance.

Can you extend this to an actor-critic method?

Certainly! Here's an extension of the previous code using the actor-critic method to solve the CartPole problem:

Message ChatGPT

ChatGPT can make mistakes. Check important info.

[ChatGPT](#)

BERT family



dmlc
XGBoost

image synthesis



Prompt: Epic anime artwork of a wizard atop a mountain at night casting a cosmic spell into the dark sky that says "Stable Diffusion 3" made out of colorful energy

[Stable Diffusion 3 — Stability AI](#)

Generative vs Predictive Models

text-to-video



Prompt: A stylish woman walks down a Tokyo street filled with warm glowing neon and animated city signage. She wears a black leather jacket, a long red dress, and black boots, and carries a black purse. She wears sunglasses and red lipstick. She... +

[Sora | OpenAI](#)

tabular data

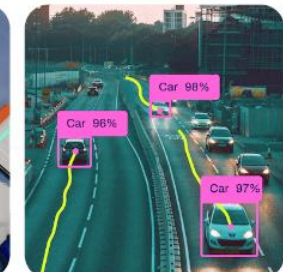
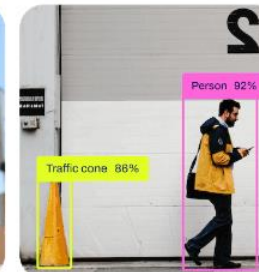
computer vision

Classify

Detect

Segment

Track



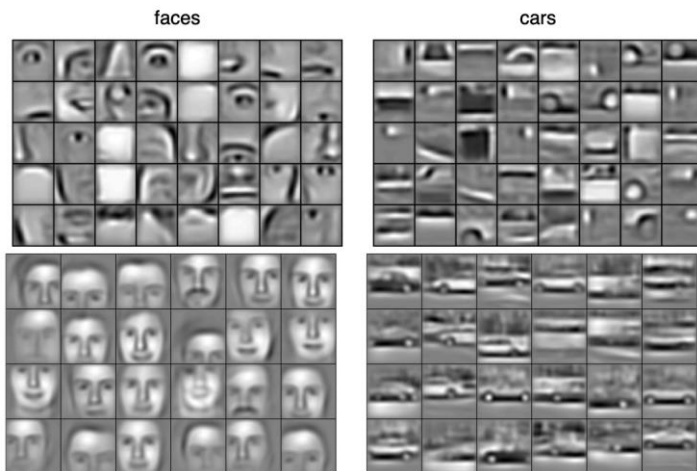
[YOLO](#)

Ladder of Generalization

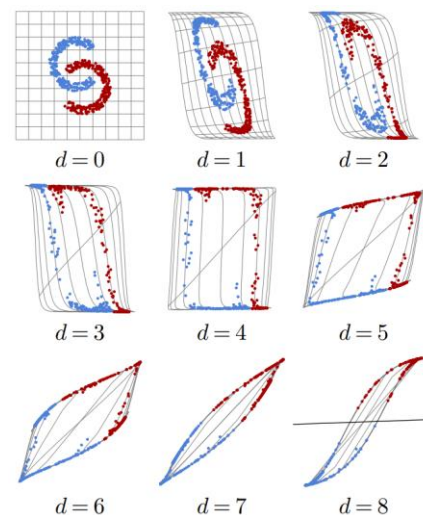
classic ML: feature engineering

deep learning: feature learning

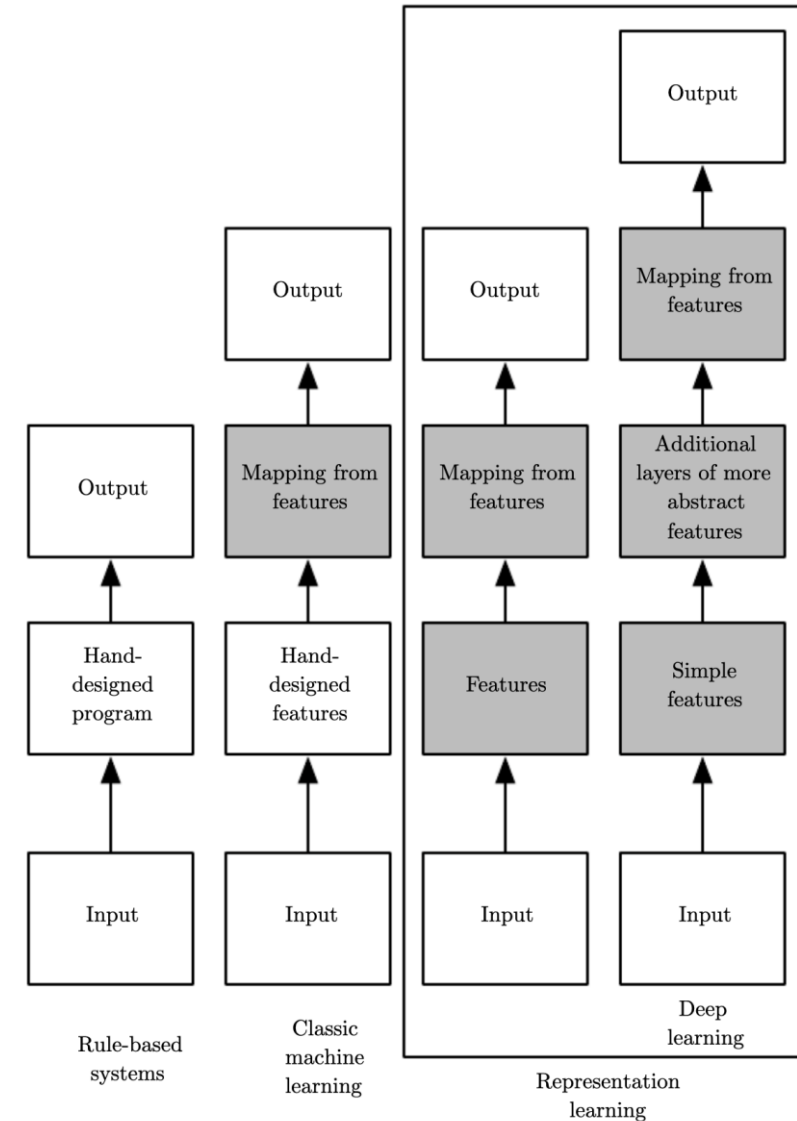
(hierarchy of concepts learned from raw data in deep graph with many layers)



[source](#)



[source](#)



[source](#)

Structured/Tabular vs Unstructured Data

unstructured data: homogenous

→ deep learning rules



ImageNet

The Lord of the Rings

Article Talk

From Wikipedia, the free encyclopedia

(Redirected from Lord of the rings)

 This article is about the book. For other uses, see The Lord of the Rings (disambiguation).
 "*War of the Ring*" redirects here. For other uses, see War of the Ring (disambiguation).

The Lord of the Rings is an epic^[1] high fantasy novel^[2] by the English author and scholar J. R. R. Tolkien. Set in *Middle-earth*, the story began as a sequel to Tolkien's 1937 children's book *The Hobbit*, but eventually developed into a much larger work. Written in stages between 1937 and 1949, *The Lord of the Rings* is one of the best-selling books ever written, with over 150 million copies sold.^[3]

The title refers to the story's main antagonist,^[4] Sauron, the Dark Lord who in an earlier age created the One Ring to rule the other Rings of Power given to Men, Dwarves, and Elves, in his campaign to conquer all of Middle-earth. From homely beginnings in the Shire, a hobbit land reminiscent of the English countryside, the story ranges across Middle-earth, following the quest to destroy the One Ring, seen mainly through the eyes of the hobbits Frodo, Sam, Merry, and Pippin. Aiding Frodo are the Wizard Gandalf, the Men Aragorn and Boromir, the Elf Legolas, and the Dwarf Gimli, who unite in order to rally the Free Peoples of Middle-earth against Sauron's armies and give Frodo a chance to destroy the One Ring in the fire of Mount Doom.

Although often mistakenly called a trilogy, the work was intended by Tolkien to be one volume in a two-volume set along with *The Silmarillion*.^{[5][6]} For economic reasons, *The Lord of the Rings* was first published over the course of a year from 29 July 1954 to 20 October 1955 in three volumes rather than one^[6] under the titles *The Fellowship of the Ring*, *The Two Towers*, and *The Return of the King*. The *Silmarillion* appeared only after the author's death. The work is divided internally into six books, two per volume, with several appendices of background material.^[7] These three volumes were later published as a boxed set, and even finally as a single volume, following the author's original intent.

structured data: heterogenous

→ feature engineering needed

→ deep learning loses its advantage over shallow methods

→ e.g., gradient boosting still prominent

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities	...	PoolArea	PoolQC	Fence	MiscFeature	MiscVal	MoSold	YrSold	SaleType	SaleCondition	SalePrice
9	1	60	RL	65.0	8450	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	NaN	NaN	0	2	2008	WD	Normal	208500
1	2	20	RL	80.0	9600	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	NaN	NaN	0	5	2007	WD	Normal	181500
2	3	60	RL	69.0	11250	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN	NaN	0	9	2008	WD	Normal	223500
3	4	70	RL	60.0	9550	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN	NaN	0	2	2006	WD	Abnorm	140000
4	5	60	RL	84.0	14260	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN	NaN	0	12	2008	WD	Normal	250000
...
1455	1456	60	RL	62.0	7917	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	NaN	NaN	0	8	2007	WD	Normal	175000
1456	1457	20	RL	85.0	13175	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	PrvPrv	NaN	0	2	2010	WD	Normal	210000
1457	1458	70	RL	66.0	9842	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	GdPrv	Shed	2500	5	2010	WD	Normal	265500
1458	1459	20	RL	68.0	9717	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	NaN	NaN	0	4	2010	WD	Normal	142125
1459	1460	20	RL	75.0	9937	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	NaN	NaN	0	6	2008	WD	Normal	147500

[1460 rows x 21 columns]

Transfer Learning

idea:

- generic pre-training of foundation models on huge data sets
- subsequent fine-tuning for specific tasks on small(er) data sets

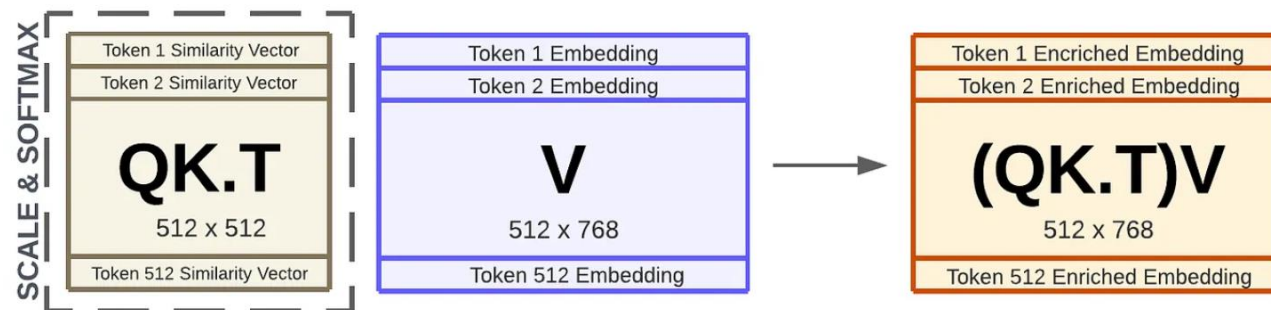
very successful for:

- computer vision (e.g., object classification)
- language models (e.g., BERT, GPT)

not (yet) for tabular data

Language Models: Contextual Semantics

- self-supervised learning: e.g., next/masked-word prediction
- tokenization: split text into chunks (e.g., words)
- semantics by means of vector embeddings: e.g., via bag-of-words (or end-to-end in transformer)
- positional encoding & embeddings: order of sequence
- contextual embeddings: (self-)attention (weighted averages: influence from other tokens)

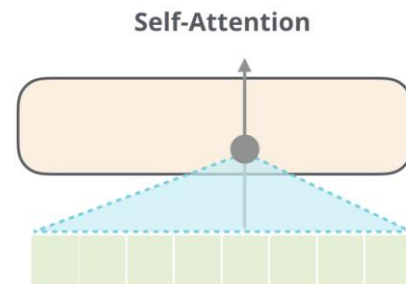


[source](#)

Encoder vs Decoder LLMs

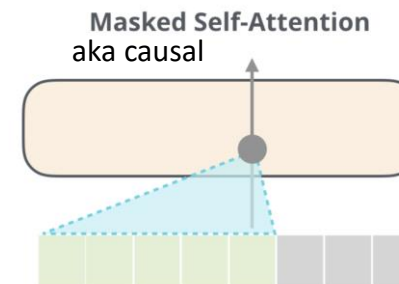
encoder-only LLMs

- prime example: BERT
- self-supervised pre-training: masked-word prediction
- fine-tuning on downstream tasks (e.g., sequence classification)
- can't generate text
- can't be prompted



decoder-only LLMs

- prime example: GPT
- self-supervised pre-training: next-word prediction
- instruction tuning (e.g., RL from human feedback)
- generate text: chat bots
- prompt engineering



Generative vs Predictive/Discriminative Models

discriminative models:

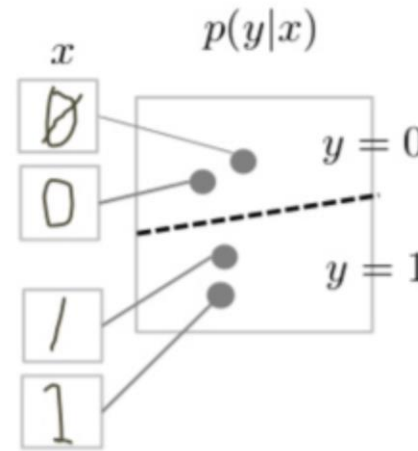
predict conditional probability $P(Y|\mathbf{X})$

generative models:

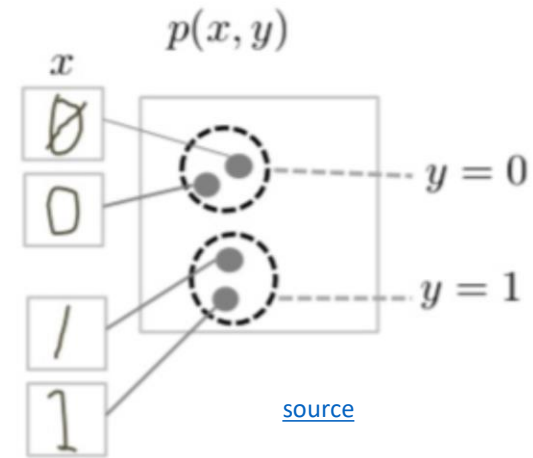
predict joint probability $P(Y, \mathbf{X})$

(or just $P(\mathbf{X}) \rightarrow$ unsupervised learning)

discriminative model



generative model



task of generative models more difficult: need to model full data distribution rather than merely find patterns in inputs to distinguish outputs

generative models allow to generate new data samples (text, images, video, proteins, ...)

predictive models usually better for predictive tasks, business problems often specific/predictive

Deep Learning for Generative AI

Depending on the application, there are currently two dominant approaches for generative AI:

- text generation: decoder LLMs
- image synthesis: diffusion models

note the difference between image synthesis and multimodal understanding in LLMs
(images as additional input sequences to transformer, tokenized by splitting into patches)


Image Synthesis

idea: generate new images as variations of training data

usually conditioned on text (prompt) by transformers

compared to text generation, additional mechanism needed (e.g., diffusion) to create more complex structures

example: DALL-E 2



Text Generation

in-context learning as alternative to fine-tuning (new paradigm):

feed information into LLM via input prompt (decoder LLMs)

→ attention to context

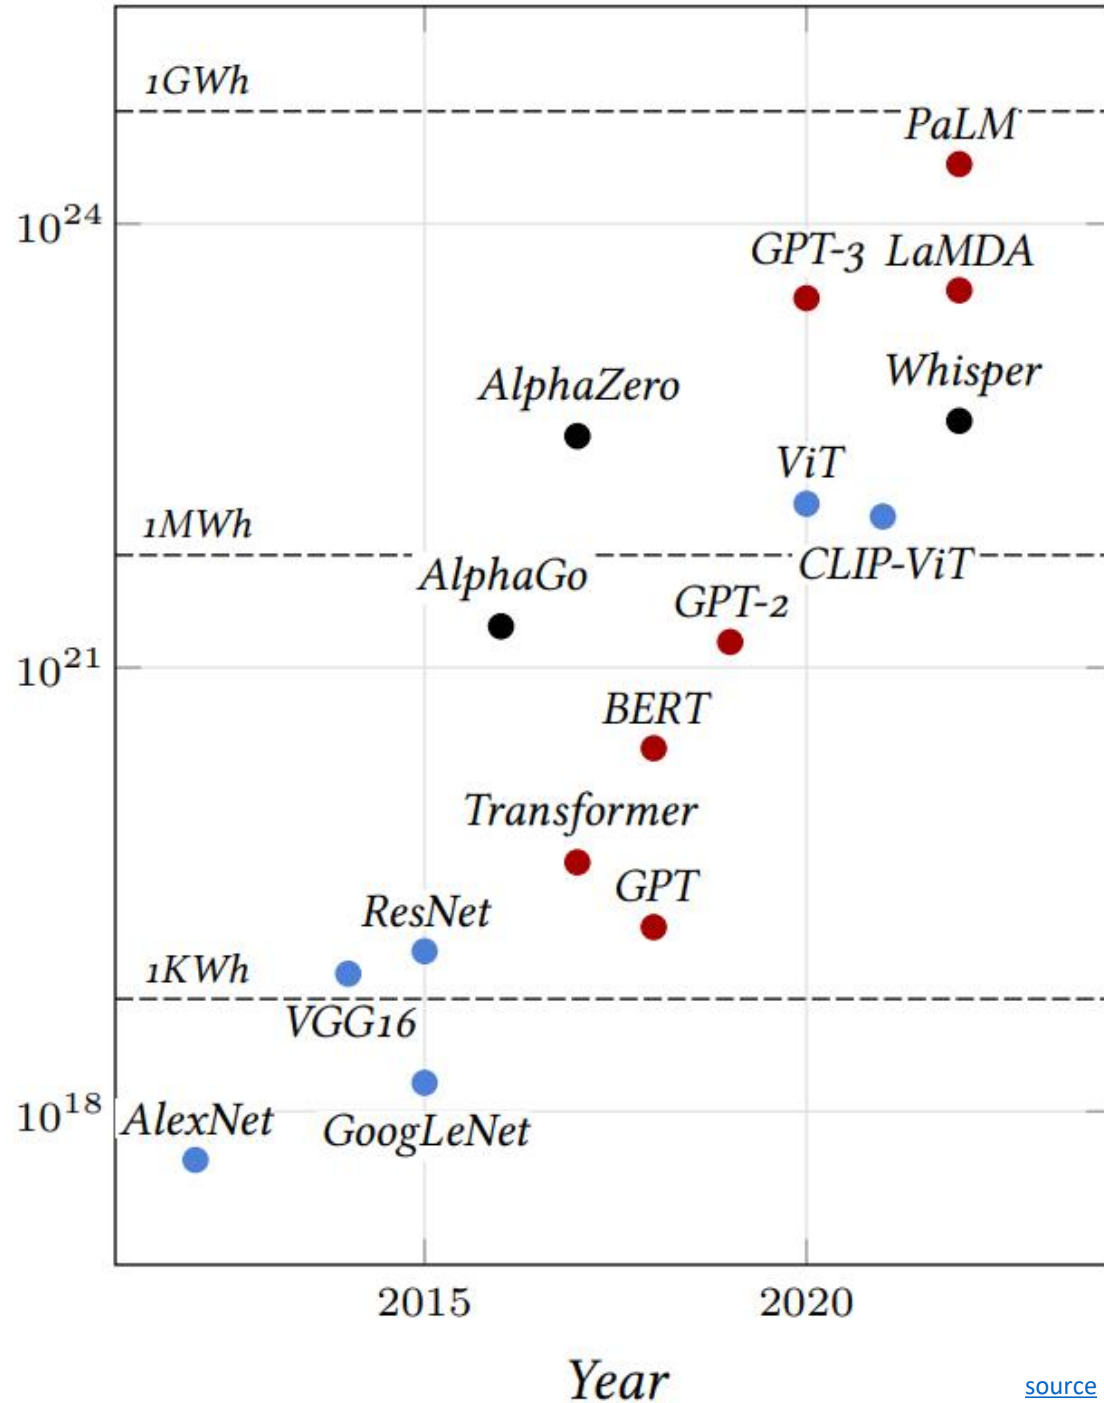
typical prompt:

instructions, context (potentially retrieved), query, output indicator

enables multi-task capabilities (all of ML before was only narrow tasks)

→ assistants

Training cost (FLOP)



[source](#)

(open-source) SOTA (July 2024):
[Llama 3](#)

Closed-source vs. open-weight models

Llama 3.1 405B closes the gap with closed-source models for the first time in history.

@maximelabonne



Some LLM Numbers

example Llama 3 405B:

- vocabulary size (tokens): 128K
 - embedding/model dimension: 16,384
 - parameters: 405B
 - training tokens: 15.6T
 - context length/window (tokens): 128K
 - training hardware: 16K GPUs (H100)
- ← factor less than 40
→ a lot of memorizing

LLM Agents

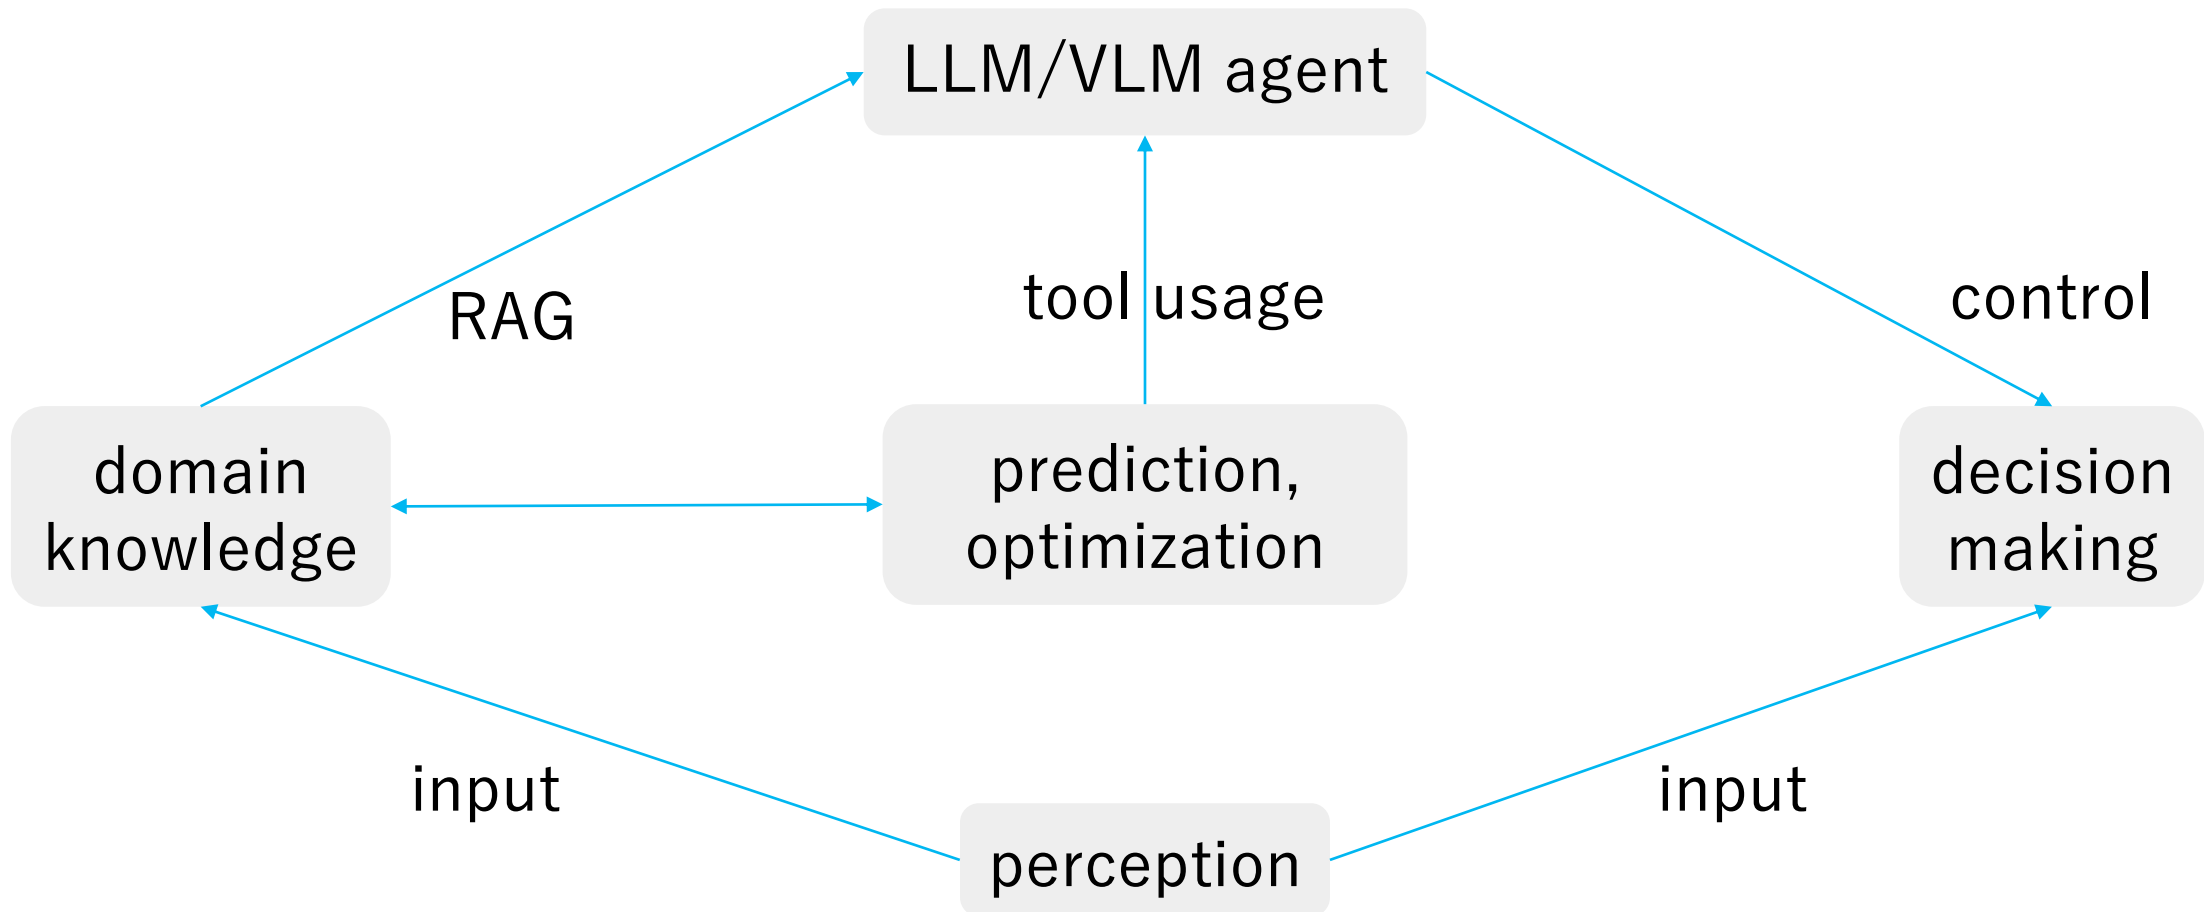
current AI good at learning statistical patterns and making predictions

but no real “understanding”, and limited reasoning and planning capabilities

desired agent capabilities:

- planning (LLM: decomposition of complex issue in multiple simple steps)
- tool use (LLM: use predictive models for numerical/optimization tasks)
- reflection
- collaboration with other agents

Goal: Autonomous End-to-End Workflow



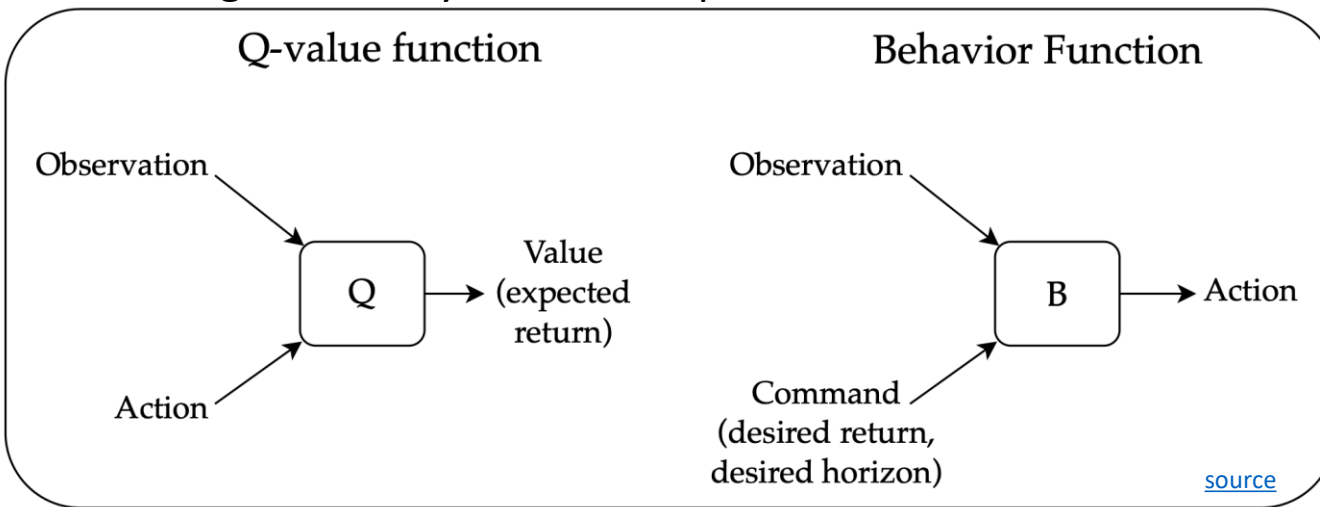
Sequential Decision Making

typically, domain of reinforcement learning

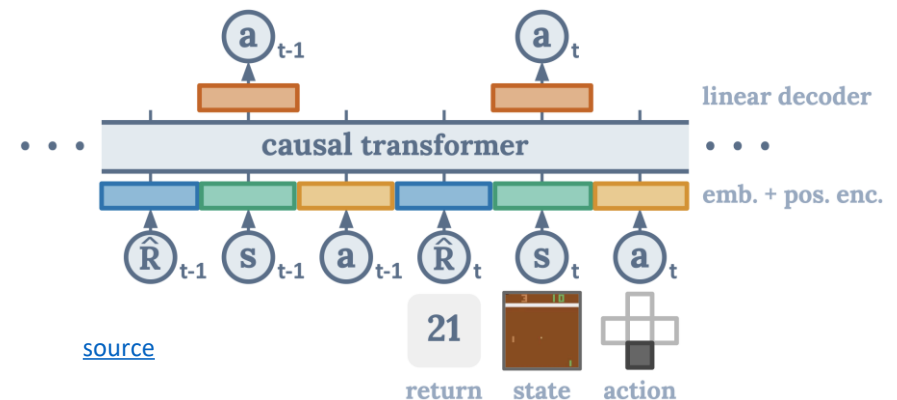
sequence modeling as alternative:

- generative: transformer decoder to autoregressively model trajectories
- credit assignment directly via self-attention: state-return associations
- desired return tokens as prompt for action generation

overcoming the deadly triad of deep RL:

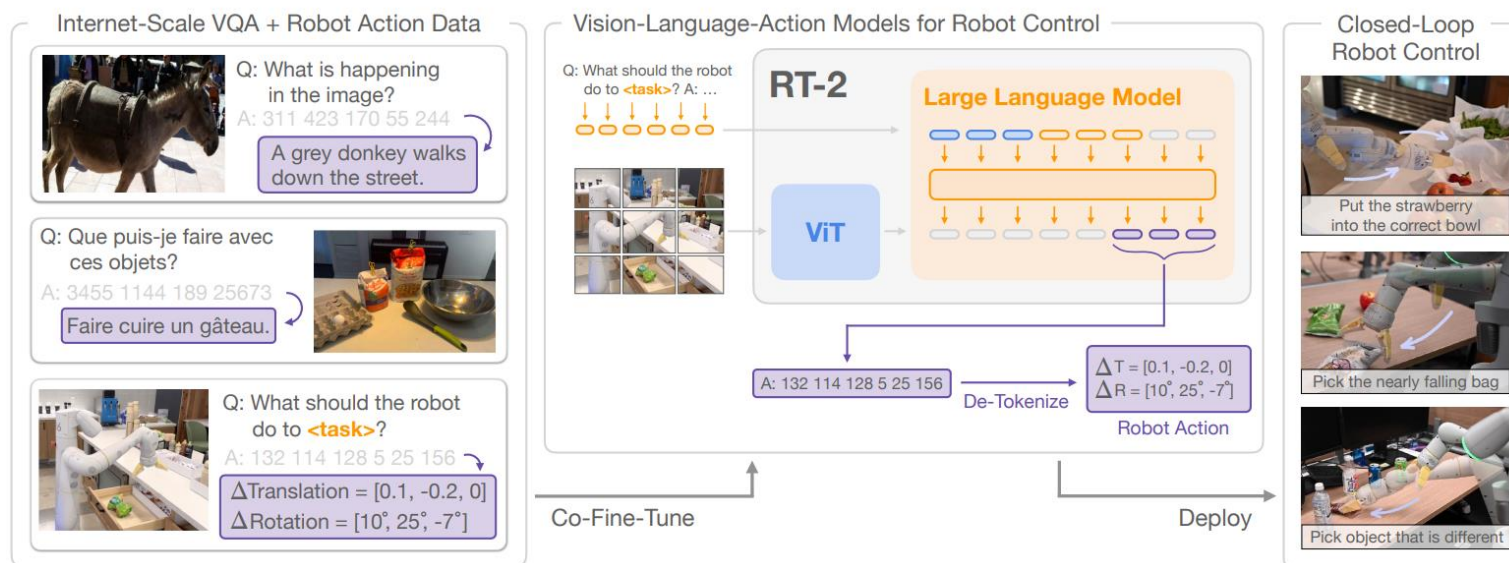


Decision Transformer:



Robotic Control generated by LLMs/VLMs

RT-2:



Code as Policies:

SayCan (grounding with pre-trained skills):

